

WhitePaper: Cyber-Security Variant of TruCol protocol

Eliminating triage intermediaries for zero-day exploits using a decentralised payout protocol

Chihab Amghane

Radboud University

Victoria Bosch

Radboud University

Rashim Charles

Radboud University

Marc Droogh

Delft University of Technology

Clara Maine

Radboud University

Chinari Subhechha Subudhi

Delft University of Technology

Akke Toeter ✉ 

Delft University of Technology, Radboud University

Eric van der Toorn

Delft University of Technology

1 Introduction

This document presents a Trustless Security protocol that aims to help ethical hackers retrieve zero-day exploit bounties without ambiguity, whilst simultaneously enabling companies to show their customers how much money is staked on their open source software stacks.

To explain how the protocol may help both of these stakeholders (ethical hackers and companies using open source software), we will first describe, what we think is, a typical procedure for vulnerability disclosures in Section 2. Then we will explain how the protocol can improve upon that in Section 3. Next, Section 4 describes strategies specifying how the protocol may be implemented. The limitations and weaknesses of our strategy and protocol are detailed in Section 5. This white-paper is concluded in Section 6.

2 Assumptions

This section starts with the assumptions that we made about the way vulnerability disclosures of zero-day exploits are currently typically handled.

2.1 Ethical Hacker Perspective

1. We assume it is not always convenient and beneficial for ethical hackers to publish an exploit and retrieve an accompanying financial reward. This assumption is based on private communications with two ethical hackers and popular media such as darknet diaries, posts on news.ycombinator.com, and possibly other sources. This assumption is based on (a combination of) the following sub-assumptions:

- 40 a. If ethical hackers give a complete disclosure of the vulnerability directly to the company,
41 they lose their negotiating leverage. A trusted triage intermediary is often used to
42 overcome this issue.
- 43 b. The effort required to contact the company and convince them of the seriousness of
44 the bug may consume unnecessary resources.
- 45 c. Cautiousness from the ethical hacker with respect to the legality of discovering the
46 vulnerability when approaching the company may hinder/slow down the vulnerability
47 disclosure process.
- 48 d. Ambiguity in the specification of the bug bounty/reward program may be interpreted
49 in the advantage of the company during triage.
- 50 e. The triage process may take a relatively long time, requiring the ethical hacker to have
51 sufficient funds to sustain living costs coverage until the pay-out.
- 52 f. The triage intermediary consumes financial resources, which lower the amount allocat-
53 able to the hacker, and/or raise the cost of cyber-security defences.
- 54 g. Vulnerabilities may be discovered at small/non-profit software development companies
55 that have not allocated a large budget fraction to security. This may render navigating
56 the vulnerability disclosure process successfully, challenging for such small/non-profit
57 organisation.

58 2.2 Company Perspective

- 59 1. We assume cybersecurity vulnerabilities become increasingly more relevant in our increas-
60 ingly more digitized world. This assumption may be seen as being substantiated by for
61 example the *Cyber Security Assessment Netherlands 2021 (CSAN 2021)* as presented by
62 the Dutch National Coordinator Counterterrorism and Safety of the Ministry of Justice
63 and Security. Currently, there is only the Dutch version available at: <https://www.nctv.nl>.
64 We assume that this trend can be extrapolated from a Dutch perspective to a more global
65 perspective, given the international media coverage of many ransomware attacks.
- 66 2. We assume that companies are interested, or will become more interested, in showing their
67 customers and/or stakeholders (a quantified perspective on) *how* secure their technology
68 is. We assume it can be quite challenging to convey this perspective clearly due to the
69 following factors:
 - 70 a. Vulnerabilities can be found in various sections of the company, ranging from social
71 engineering, misconfiguration to zero-day exploits. It is difficult to give customers a
72 comprehensive yet concise/simple insight in "how secure" each of these attack surfaces
73 are.
 - 74 b. The impact of a vulnerability may be ambiguous or not easily quantifiable. For example,
75 for some companies, vulnerabilities may allow malicious actors to take over critical
76 infrastructure, whilst other vulnerabilities may lead to data leaks or other undesired
77 side effects.
 - 78 c. It may be difficult to accurately assess the capabilities of malicious adversaries.
- 79 3. We assume some companies might be unfamiliar with vulnerability disclosure and ac-
80 companying triage processes. These delicate processes may seem intimidating for new
81 companies that want to start paying attention to their cybersecurity, and this may lead
82 to a lower allocation of cybersecurity budget.

3 Protocol

This section presents the TruSec protocol, and explains how it can improve the way vulnerability disclosures are completed for deterministically verifiable zero-day exploits.

3.1 Scope

The protocol is primarily designed to automate deterministically verifiable zero-day exploits vulnerability disclosures. It can also be used to hedge against misconfigurations and supply-chain attacks. For example, companies can add a specific configuration (yaml) to the DVM, and add a bounty on that forked DVM. This way, a hacker may leverage the particular configuration to find an exploit. This procedure also allows the protocol to identify some supply-chain attack vulnerabilities. For example, if an invalid certificate is used to compromise the device.

However, both misconfiguration and supply chain attack partially deviate from the main benefit of collective nature of the protocol. For example, it may incentivise hackers to focus efforts on particular configurations, that are not (necessarily) useful for other companies. However, at the same time, hackers could still opt to focus on the mutual elements of all forked decentralised virtual machines (DVMs) to collect the bounties with a single, more powerful exploit. This scope/applicability of the protocol is visualised in Figure 1.

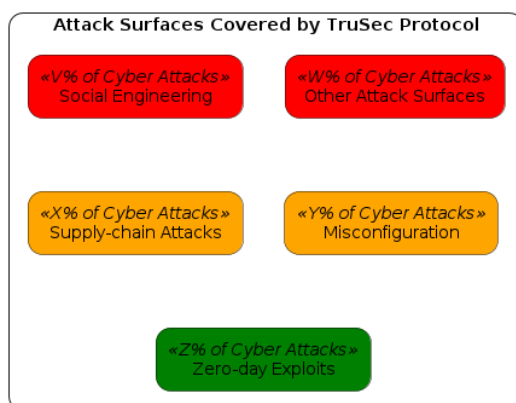


Figure 1 The proposed TruSec protocol is not suited to deal with social engineering attacks, nor is it ideal for misconfiguration exploits and/or supply-chain attacks. Instead, it is designed to increase the rate of discovery of deterministically verifiable zero-day exploits. Note, we acknowledge that attacks can be, and often are, a combination of the types.

With respect to Figure 1, the following notes are made:

1. The orange attack types imply the proposed protocol is not designed to tackle these issues, nor does it provide full coverage (against malicious agents) for these attack types. However:
 - a. The misconfiguration could be covered if companies upload their configurations into DVMs. These configurations would typically not benefit from the collaborative staking, as it is less likely that other companies happen to use the same configurations.
 - b. Some of the supply chain attacks could be covered if the ethical hackers are able to propagate these supply chain exploits into the DVMs.

109 3.2 Usage

110 With this scope defined, one can look at how companies and ethical hackers interact according
111 to the proposed protocol.

112 The basic idea is that companies and users (stakeholders) can put their open source software
113 stacks on a decentralised virtual machine DVM. They can then collectively stake money
114 on the security of the stacks, such that everyone can see how much money says: *the use*
115 *of certain software packages/combinations is safe*. This enables companies, to show their
116 customers for example:

117 *With us, your data is stored using MongoDB Version 5.1, \$314.159,- says it is uncompromised,*
118 *and it's running on Ubuntu Server version 21.10, which has \$4.200.000,- staked on its security.*
119 *This setup has a configuration with a security on which we staked \$9001,-. If any of these*
120 *software packages get compromised by ethical hackers, we will be the first to know.*

121 We believe that might be clear language that enables decision makers and customers interested
122 in company A, to get a simple, and intuitive understanding on: *how secure* is the software
123 stack of company A?

124 For the ethical/ethical hackers, the advantages are clear; they know before they start their
125 work how large their payout will be, and they get a direct payout upon completion of their
126 work (after the predetermined responsible disclosure period has ended).

127 3.2.1 Disclaimer

128 The presented protocol does not provide insight in the complete security of a system/company.
129 As visualised in Figure 1, the protocol does not cover all attack surfaces of companies. Hence, if
130 other attack surfaces, such as social engineering are used, companies can still get compromised,
131 regardless of the amount they staked. Therefore, it is important that the numerical value
132 of the amount staked on the zero-day exploit security level is not abused to convey a false
133 sense of security by the staking companies to their customers.

134 3.3 Description

135 The protocol is shown in Figure 2.

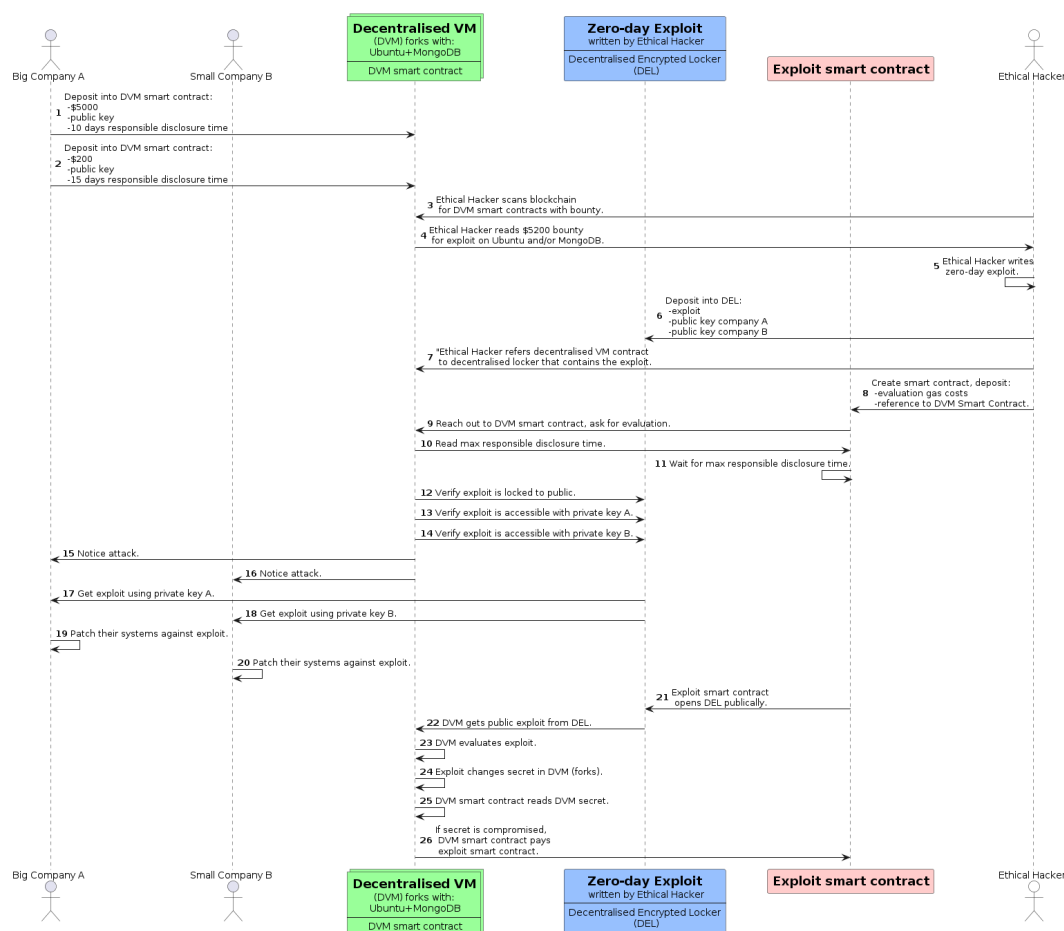


Figure 2 Visualisation of the interaction of the TruSec protocol. This is an ever-lasting cycle, where at the end of the process, companies can re-deploy the patched decentralised stack, and allocate new funds. ethical hackers can scan for new attacks.

To summarize, the protocol enables companies and users of open source software to collectively stake money on the software stacks of their choice. Hackers can see these bounties and write a new zero-day exploit for these staked systems. Next, they can deploy their attack to a decentralised locker that only opens to the stakeholders. Only the hacker and the stakeholders can then see the zero-day exploit within the responsible disclosure time (RDT) they specified. Stakeholders can then patch their systems. After the RDT, the exploit is evaluated, and if it compromises the software stack, the hacker automatically receives the staked bounty. The cycle can then start over, with companies re-deploying their patched systems and applying new stakes to their respective security.

3.3.1 Figure 2 notes

With respect to Figure 2, the following notes are made:

1. The attack written by the ethical hacker should be accessible on chain, such that everyone can verify that the attack indeed compromises the decentralised VM/honeypot. This is critical for the automatic payout.
2. The decentralised locker is used to prevent malicious hackers to inspect/copy the attack

151 before the responsible disclosure period is over.

152 3. It would be better if the DVM smart contract specifies the locker location, while granting
 153 hackers write-access, and read-access only to (certified) stakeholders. This would prevent
 154 other hackers from knowing there is a vulnerability discovered in a certain software stack
 155 before the responsible disclosure period is over. We expect such a signal might attract
 156 unwanted attention. However, at the time of writing, no mechanism is designed that
 157 would prevent people from determining whether a hacker has deployed an attack, even
 158 if it is encrypted, whilst still allowing stakeholders to access/read the attack within the
 159 responsible disclosure. This weakness is discussed in Section 5.

160 3.4 Incentives

161 To convey a better understand of the protocol, some of the incentives are evaluated along with
 162 the relevant steps shown in Figure 2. Starting with the RDT, companies and stakeholders
 163 need to have sufficient time to patch their systems, that is why they have the liberty to
 164 specify the disclosure time they need. Since the attack needs to be private until the maximum
 165 specified RDT has passed, a check is performed to verify it is indeed hidden to the public in
 166 step 12 of Figure 2.

167 3.4.1 Malicious staking

168 Malicious actors could stake a minimum amount to set the maximum RDT to an unreasonable
 169 amount, such as 3 centuries. To mitigate this tampering, hackers can select which bounties
 170 they want to collect. This allows them to forfeit the trivial stake with unreasonable RDT.

171 Another act of malicious staking could be to have a malicious agent stake to get access
 172 to the attack to abuse it within the RDT. In the purest decentralised form of the protocol
 173 there is no defence against this. To alleviate this concern, different strategies can be pursued.
 174 For example, limited access to the zero-day exploits can be granted to a ring of trusted
 175 companies using self-sovereign identity. This would however introduce subjectivity into the
 176 protocol, and it would lower the decentralised nature of the protocol, even when it is done in
 177 the form of a DAO. Alternatively, access to the exploit within RDT could be granted to the
 178 highest staker only, whilst requiring stakers that want early access to identify themselves
 179 using SSI. This would still allow malicious actors to "buy" the exploits within the RDT, like
 180 they can currently. However, it would at least inform relevant companies of their exposed
 181 position and allows them to outbid the malicious actors.

182 3.4.2 Malicious Exploit Publications

183 Hackers could also try to sabotage the protocol. For example, they could submit invalid
 184 solutions to clog the network. To prevent this, step 8 requires the hacker to deposit evaluation
 185 costs in order to be eligible for a payout.

186 Another attack from a hacker perspective could be to deny access to the stakers. This
 187 is why the companies should present a public key in steps 1 and 2, why the hacker should
 188 provide access for these keys in step 6, and the DVM smart contract should verify this access
 189 in steps 13 and 14.

190 Another form of malicious exploit publication would be to claim the bounty using the
 191 protocol and then to publish the exploit somewhere else anonymously within the RDT. The
 192 purely decentralised form has no defence against this behaviour. If it occurs frequently, a
 193 DAO/voting mechanism could be implemented that determines whether the exploit has

remained hidden to the public, however this introduces subjectivity into the protocol and opens new kinds of attack surfaces.

3.5 Added Value

The protocol allows companies to show their users how secure their open sources software stacks are against deterministically verifiable zero-day exploits. This can be done by showing the users how much money is "staked" on the security of their respective systems. This simplifies the comparison that customers can make between the security of companies against deterministically verifiable zero-day exploits. Additionally, companies (as well as the users of these open source software stacks) can re-adjust their funds and resource allocations regarding cyber-security based on this insight. This mechanism could, over multiple cycles, result in a more predictable zero-day exploit landscape.

To improve the ease of use and practical application of the protocol, a variant could be written that allows for staking on decentralised containers instead of decentralised virtual machines. This could reduce the required computational resources and broaden the usecases of the protocol to containerized applications (with minimal attack surfaces). Additionally, using hardware emulators such as QEMU, in decentralised format, companies could use this protocol to hedge against certain types of hardware exploits as well. We hope this allows extending the TruSec protocol to show consumers *"how secure"* their phones are, in real-time.

4 Implementation

Implementation details are omitted at this stage. One can note that developing decentralised virtual machines for security purposes requires a significant effort, even when considering ports from e.g. the Ethereum Virtual Machine (EVM).

5 Discussion

The presented proposal for protocol development can be critically evaluated. This section aims to identify possible weak points.

5.1 Limitations

The following limitations are identified in the protocol:

1. The proposed protocol, in its initial form, does not (necessarily) work for security compromises that are not clearly pre-defined. For example, if the decentralised virtual machine/stack/honeypot is configured to only pay-out in case an internal value/secret is modified, a ethical hacker might be able to gain read-access to the secret, which could be considered a hack, but the ethical hacker would not receive a payout. Accordingly, companies may specify different payouts to different types of security breaches. This may reduce the added value of collaborative staking.
2. The running a decentralised virtual machine with a high degree of decentralisation, along with their interactions is expected to be costly. This expectation is based on approximate costs of roughly 50 dollars for a single Ethereum transaction.
3. We expect most hacks do not rely on pure zero-day exploits. Accordingly we think the scope of this protocol is significantly limited w.r.t. the complete cybersecurity threat landscape.

- 235 4. This protocol will most likely not allow companies to test their entire system, as we
 236 currently consider it practically infeasible to simulate the various types of social engineering
 237 and or interactions with non-decentralised platforms (on a blockchain). So companies
 238 cannot, in good conscience, make claims about their overall level of cybersecurity based
 239 on this protocol alone.
- 240 5. This protocol does not protect against economically irrational malicious agents. Examples
 241 could be:
- 242 a. Actors with revenge sentiment. They could for example skip the payout and use
 243 zero-day exploits to hurt a company that staked their open source software stack.
 - 244 b. Nation states may not care about payouts and instead use found zero-day exploits
 245 themselves, instead of disclosing them.

246 5.2 Related Work

247 It was noted during the TechEx conference, that companies like Google and Microsoft already
 248 fund vulnerability disclosures for, for example, Ubuntu. This can be seen as collective funding,
 249 hence one could argue the added value of the proposed protocol may be limited in this
 250 respect.

251 Additionally, there are companies like HackerOne that perform independent triage, hence
 252 one could argue the added value of doing this in a decentralised fashion is limited.

253 6 Conclusion and Recommendations

254 The proposed protocol can enable companies to convey a quantitative level of security of
 255 (segments of) their open source technology stack to their customers. Customers can use
 256 this information on the minimum price for zero-day exploits, and compare it to their costs
 257 of suffering from a malicious zero-day exploit. This may allow them to (re)allocate their
 258 funds and cyber-security resources based on the accompanying risk-profile that emerges after
 259 multiple patching cycles with the TruSec protocol.

260 Additionally, the protocol enables ethical hackers to retrieve payouts directly without
 261 ambiguity.

262 Since computational budgets typically are costly on decentralised computing platforms,
 263 a recommendation is included to investigate the option to allow staking decentralised
 264 containerized applications instead of complete decentralised virtual machines.